Nguyen Viet Anh

Data Engineer | Applied Machine Learning | Ho Chi Minh City
(+84) 396 319 888 | nguyenvietanh1803.hcmut@gmail.com
github.com/VietAnh1803 | LinkedIn | Portfolio

Summary

Data Engineer at OSAS via AI-MED, with a Computer Science degree from HCMUT. Build Oracle ETL pipelines and Python applications, and investigate reporting errors at the source. Project work includes testing models on future transactions and evaluating stress and OCR models for a wearable-health app.

Education

Ho Chi Minh City University of Technology (HCMUT) 2022 – 2026
Bachelor of Computer Science, Information Systems – Graduated

Experience

Data Engineer – OSAS Joint Stock Company via AI-MED Jun 2025 – Present

- Develop Oracle warehouse loads from three source systems using PL/SQL and ODI 12c, including incremental loading and SCD Type 2 dimensions. Configured GoldenGate CDC for a separate public-sector warehouse.
- Reconciled source and warehouse records; corrected sales-channel classifications across approximately 68,000 rows and a policy-counting error that returned 82 policies instead of one for a customer.
- Built a Vietnamese financial-statement assistant with Python, FastAPI, and Ollama: section summaries, rule-based figure extraction, financial ratios, and source-page references. Implemented SQL/vector retrieval and an ingestion queue.

Projects

LeakageLab – Temporal Risk Benchmark Oct 2026
Independent project; synthetic data | GitHub | Kaggle

- Built a 50,000-transaction synthetic benchmark with transaction histories and seven-day delayed labels. Enforced time-based splits and wrote an ID-matched submission grader; all 18 tests passed locally and on Kaggle.
- Selected logistic regression over dummy and gradient-boosting baselines on validation. Final AP was 0.193 on 10,051 future transactions (11.82% positive); removing history features reduced validation AP from 0.132 to 0.088.

HealthMate – Wearable Health Monitoring Sep 2025 – May 2026
Graduation thesis; team of three | Backend | Flutter app

- Designed the Flutter interface and helped evaluate stress models using 869 labeled windows from 15 WESAD participants. Held out one participant at a time to assess performance on people outside the training set.
- Evaluated prescription OCR on 66 manually labeled samples. PaddleOCR with OpenCV and extraction rules achieved medication-name micro-F1 of 0.46, compared with 0.24 for PaddleOCR-VL; used the comparison to guide the app's OCR choice.

Vietnam Stock Data Pipeline 2025 – 2026
GitHub – VNStock, Kafka, PySpark Structured Streaming, TimescaleDB, Snowflake; incremental sync with a persisted watermark, Bronze/Silver/DW modeling, SQL MERGE and Streams & Tasks. Streamlit/Plotly dashboards; Docker Compose, Nginx, ZooKeeper and Linux/systemd operations.

Technical Skills

Languages & ML: Python, SQL, PL/SQL, Dart; NumPy, Pandas, scikit-learn, CatBoost.
Streaming & ETL: Kafka/kafka-python, Spark/PySpark, VNStock, ODI, GoldenGate.
Databases: Oracle, PostgreSQL/TimescaleDB, pgvector, Snowflake (Streams & Tasks).
Apps & BI: FastAPI, Flutter, Streamlit, Plotly; psycopg2, JDBC, Snowflake Connector.
Operations: Git, Linux/Bash/systemd, Docker Compose, Nginx, ZooKeeper, Kafka UI, Spark UI/History Server, pgAdmin.

Certifications

View all certificates

Python Project for Data Engineering – IBM/Coursera, 2026.
Relational Database Administration (DBA) – IBM/Coursera, 2026.